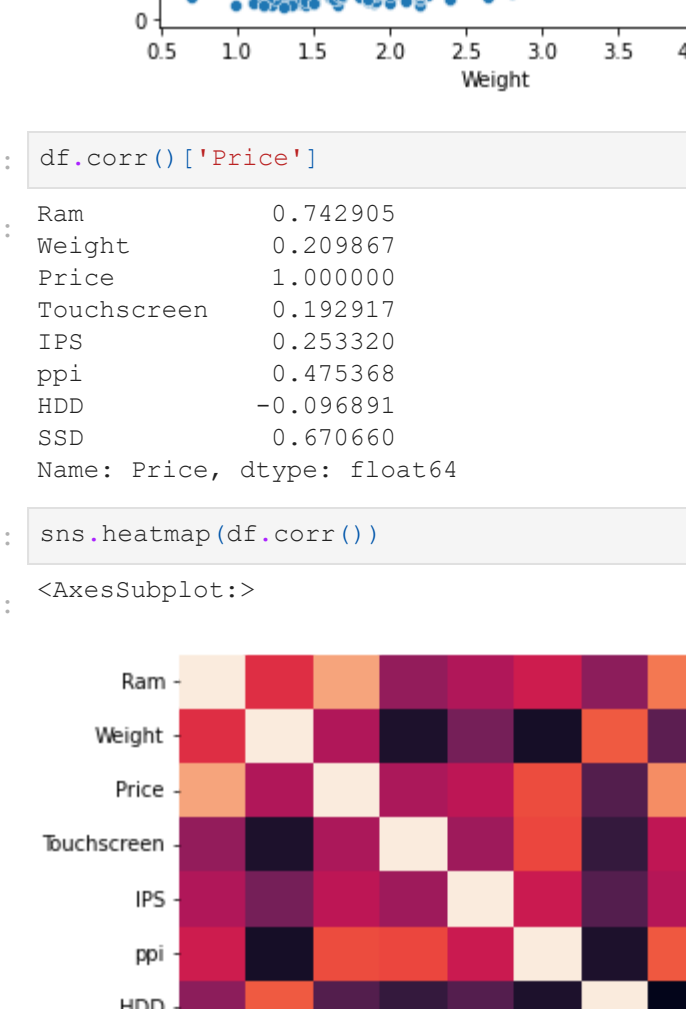
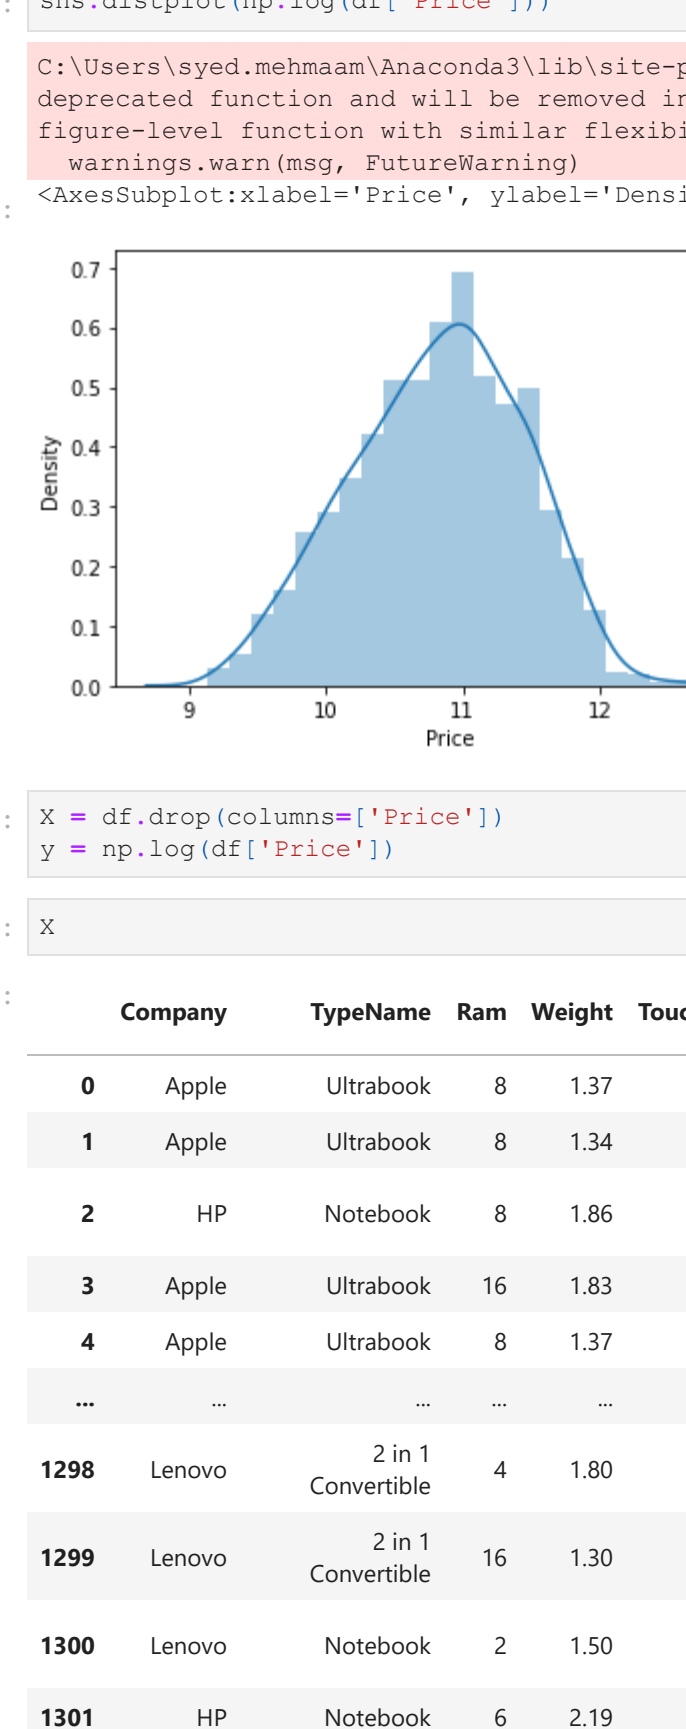


```
In [66]: sns.scatterplot(x=df['Weight'], y=df['Price'])  
Out[66]: <AxesSubplot: xlabel='Weight', ylabel='Price'>
```

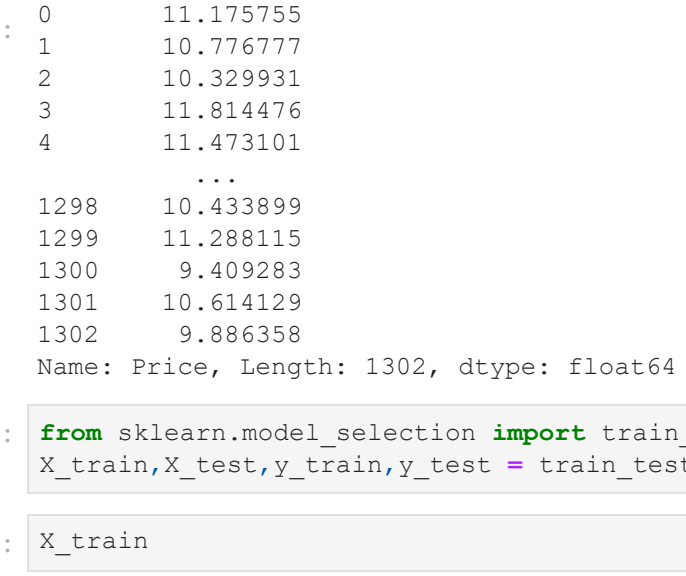


```
In [67]: df.corr()['Price']  
Out[67]:  
Ram          0.742905  
Weight       0.209867  
Price        1.000000  
Touchscreen  0.192917  
IPS          0.253320  
ppi          0.475368  
HDD          -0.096891  
SSD          0.670660  
Name: Price, dtype: float64
```

```
In [68]: sns.heatmap(df.corr())  
Out[68]: <AxesSubplot: >
```



```
In [69]: sns.distplot(np.log(df['Price']))  
C:\Users\ayed.mehman\Anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: 'distplot' is a deprecated function and will be removed in a future version. Please adapt your code to use either 'displot' (a figure-level function with similar flexibility) or 'histplot' (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)  
Out[69]: <AxesSubplot: xlabel='Price', ylabel='Density'>
```



```
In [70]: X = df.drop(columns=['Price'])  
y = np.log(df['Price'])
```

```
In [71]: X  
Out[71]:
```

	Company	TypeName	Ram	Weight	Touchscreen	IPS	ppi	Cpu brand	HDD	SSD	Gpu brand	os
0	Apple	Ultrabook	8	1.37	0	1	226.983005	Intel Core i5	0	128	Intel	Mac
1	Apple	Ultrabook	8	1.34	0	0	127.677940	Intel Core i5	0	0	Intel	Mac
2	HP	Notebook	8	1.86	0	0	141.211998	Intel Core i5	0	256	Intel	Others/No OS/Linux
3	Apple	Ultrabook	16	1.83	0	1	220.534624	Intel Core i7	0	512	AMD	Mac
4	Apple	Ultrabook	8	1.37	0	1	226.983005	Intel Core i5	0	256	Intel	Mac
...
1298	Lenovo	2 in 1 Convertible	4	1.80	1	1	157.350512	Intel Core i7	0	128	Intel	Windows
1299	Lenovo	2 in 1 Convertible	16	1.30	1	1	276.053530	Intel Core i7	0	512	Intel	Windows
1300	Lenovo	Notebook	2	1.50	0	0	111.935204	Other Intel Processor	0	0	Intel	Windows
1301	HP	Notebook	6	2.19	0	0	100.454670	Intel Core i7	1000	0	AMD	Windows
1302	Asus	Notebook	4	2.20	0	0	100.454670	Other Intel Processor	500	0	Intel	Windows

1302 rows x 12 columns

```
In [72]: y  
Out[72]:  
0      11.175755  
1      10.776777  
2      10.329931  
3      11.814476  
4      11.473101  
...
```

```
1298    10.433899  
1299    11.288115  
1300     9.409283  
1301    10.614129  
1302     9.886358  
Name: Price, Length: 1302, dtype: float64
```

```
In [73]: from sklearn.model_selection import train_test_split  
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.15,random_state=2)
```

```
In [74]: X_train  
Out[74]:
```

	Company	TypeName	Ram	Weight	Touchscreen	IPS	ppi	Cpu brand	HDD	SSD	Gpu brand	os
183	Toshiba	Notebook	8	2.00	0	0	100.454670	Intel Core i5	0	128	Intel	Windows
1141	Moshi	Gaming	8	2.40	0	0	141.211998	Intel Core i3	1000	128	Nvidia	Windows
1049	Asus	Netbook	4	1.20	0	0	135.094211	Other Intel Processor	0	0	Intel	Others/No OS/Linux
1020	Dell	2 in 1 Convertible	4	2.08	1	1	141.211998	Intel Core i3	1000	0	Intel	Windows
878	Dell	Ultrabook	4	2.18	0	0	141.211998	Intel Core i3	1000	128	Nvidia	Windows
...
466	Acer	Notebook	4	2.20	0	0	100.454670	Intel Core i5	500	0	Nvidia	Windows
299	Acer	Ultrabook	16	1.63	0	0	141.211998	Intel Core i7	0	512	Nvidia	Windows
493	Asus	Notebook	8	2.20	0	0	100.454670	AMD Processor	1000	0	AMD	Windows
527	Lenovo	Notebook	8	2.20	0	0	100.454670	Intel Core i3	2000	0	Nvidia	Others/No OS/Linux
1193	Apple	Ultrabook	8	0.92	0	1	226.415547	Other Intel Processor	0	0	Intel	Mac

1106 rows x 12 columns

```
In [75]: from sklearn.compose import ColumnTransformer  
from sklearn.pipeline import Pipeline  
from sklearn.preprocessing import OneHotEncoder  
from sklearn.metrics import r2_score,mean_absolute_error
```

```
In [76]: from sklearn.linear_model import LinearRegression,Ridge,Lasso  
from sklearn.neighbors import KNeighborsRegressor  
from sklearn.tree import DecisionTreeRegressor  
from sklearn.ensemble import RandomForestRegressor,GradientBoostingRegressor,AdaBoostRegressor,ExtraTreesRegressor  
from sklearn.svm import SVR  
from xgboost import XGBRegressor
```

```
In [77]: df.head(5)  
Out[77]:
```

	Company	TypeName	Ram	Weight	Price	Touchscreen	IPS	ppi	Cpu brand	HDD	SSD	Gpu brand	os
0	Apple	Ultrabook	8	1.37	71378.6832	0	1	226.983005	Intel Core i5	0	128	Intel	Mac
1	Apple	Ultrabook	8	1.34	47895.5232	0	0	127.677940	Intel Core i5	0	0	Intel	Mac
2	HP	Notebook	8	1.86	30636.0000	0	0	141.211998	Intel Core i5	0	256	Intel	Others/No OS/Linux
3	Apple	Ultrabook	16	1.83	135195.3360	0	1	220.534624	Intel Core i7	0	512	AMD	Mac
4	Apple	Ultrabook	8	1.37	96095.8080	0	1	226.983005	Intel Core i5	0	256	Intel	Mac

```
In [78]: #Linear Regression  
step1 = ColumnTransformer(transformers=[  
    ('col_tnf',OneHotEncoder(sparse=False,drop='first')), [0,1,7,10,11])  
]),remainder='passthrough')  
step2 = LinearRegression()  
pipe = Pipeline([  
    ('step1',step1),  
    ('step2',step2)  
])  
pipe.fit(X_train,y_train)  
y_pred = pipe.predict(X_test)  
print('R2 score',r2_score(y_test,y_pred))  
print('MAE',mean_absolute_error(y_test,y_pred))  
R2 score 0.8073277448418734  
MAE 0.21017827976428724
```

```
In [79]: #Ridge Regression  
step1 = ColumnTransformer(transformers=[  
    ('col_tnf',OneHotEncoder(sparse=False,drop='first')), [0,1,7,10,11])  
]),remainder='passthrough')  
step2 = Ridge(alpha=10)  
pipe = Pipeline([  
    ('step1',step1),  
    ('step2',step2)  
])  
pipe.fit(X_train,y_train)  
y_pred = pipe.predict(X_test)  
print('R2 score',r2_score(y_test,y_pred))  
print('MAE',mean_absolute_error(y_test,y_pred))  
R2 score 0.8127331031311809  
MAE 0.20926802242882865
```

```
In [80]: #Lasso Regression  
step1 = ColumnTransformer(transformers=[  
    ('col_tnf',OneHotEncoder(sparse=False,drop='first')), [0,1,7,10,11])  
]),remainder='passthrough')  
step2 = Lasso(alpha=0.001)  
pipe = Pipeline([  
    ('step1',step1),  
    ('step2',step2)  
])  
pipe.fit(X_train,y_train)  
y_pred = pipe.predict(X_test)  
print('R2 score',r2_score(y_test,y_pred))  
print('MAE',mean_absolute_error(y_test,y_pred))  
R2 score 0.8071853945317105  
MAE 0.21114361613472565
```

```
In [81]: #KNN  
step1 = ColumnTransformer(transformers=[  
    ('col_tnf',OneHotEncoder(sparse=False,drop='first')), [0,1,7,10,11])  
]),remainder='passthrough')  
step2 = KNeighborsRegressor(n_neighbors=3)  
pipe = Pipeline([  
    ('step1',step1),  
    ('step2',step2)  
])  
pipe.fit(X_train,y_train)  
y_pred = pipe.predict(X_test)  
print('R2 score',r2_score(y_test,y_pred))  
print('MAE',mean_absolute_error(y_test,y_pred))  
R2 score 0.8021984604448553  
MAE 0.19319787621521116
```

```
In [82]: #Decision Tree  
step1 = ColumnTransformer(transformers=[  
    ('col_tnf',OneHotEncoder(sparse=False,drop='first')), [0,1,7,10,11])  
]),remainder='passthrough')  
step2 = DecisionTreeRegressor(max_depth=8)  
pipe = Pipeline([  
    ('step1',step1),  
    ('step2',step2)  
])  
pipe.fit(X_train,y_train)  
y_pred = pipe.predict(X_test)  
print('R2 score',r2_score(y_test,y_pred))  
print('MAE',mean_absolute_error(y_test,y_pred))  
R2 score 0.8407719845727799  
MAE 0.18296991310352723
```

```
In [83]: #SVR  
step1 = ColumnTransformer(transformers=[  
    ('col_tnf',OneHotEncoder(sparse=False,drop='first')), [0,1,7,10,11])  
]),remainder='passthrough')  
step2 = SVR(kernel='rbf',C=10000,epsilon=0.1)  
pipe = Pipeline([  
    ('step1',step1),  
    ('step2',step2)  
])  
pipe.fit(X_train,y_train)  
y_pred = pipe.predict(X_test)  
print('R2 score',r2_score(y_test,y_pred))  
print('MAE',mean_absolute_error(y_test,y_pred))  
R2 score 0.8081180902257614  
MAE 0.202390594927481507
```

```
In [84]: #Random Forest  
step1 = ColumnTransformer(transformers=[  
    ('col_tnf',OneHotEncoder(sparse=False,drop='first')), [0,1,7,10,11])  
]),remainder='passthrough')  
step2 = RandomForestRegressor(n_estimators=100,  
                             random_state=3,  
                             max_samples=0.5,  
                             max_features=0.75,  
                             max_depth=15,  
                             bootstrap=True)  
pipe = Pipeline([  
    ('step1',step1),  
    ('step2',step2)  
])  
pipe.fit(X_train,y_train)  
y_pred = pipe.predict(X_test)  
print('R2 score',r2_score(y_test,y_pred))  
print('MAE',mean_absolute_error(y_test,y_pred))  
R2 score 0.887340237832468  
MAE 0.15860130110457718
```

```
In [85]: #ExtraTrees  
step1 = ColumnTransformer(transformers=[  
    ('col_tnf',OneHotEncoder(sparse=False,drop='first')), [0,1,7,10,11])  
]),remainder='passthrough')  
step2 = ExtraTreesRegressor(n_estimators=100,  
                             random_state=3,  
                             max_samples=0.5,  
                             max_features=0.75,  
                             max_depth=15,  
                             bootstrap=True)  
pipe = Pipeline([  
    ('step1',step1),  
    ('step2',step2)  
])  
pipe.fit(X_train,y_train)  
y_pred = pipe.predict(X_test)  
print('R2 score',r2_score(y_test,y_pred))  
print('MAE',mean_absolute_error(y_test,y_pred))  
R2 score 0.8850720167552375  
MAE 0.16154538000217084
```

```
In [86]: #AdaBoost  
step1 = ColumnTransformer(transformers=[  
    ('col_tnf',OneHotEncoder(sparse=False,drop='first')), [0,1,7,10,11])  
]),remainder='passthrough')  
step2 = AdaBoostRegressor(n_estimators=15,learning_rate=1.0)  
pipe = Pipeline([  
    ('step1',step1),  
    ('step2',step2)  
])  
pipe.fit(X_train,y_train)  
y_pred = pipe.predict(X_test)  
print('R2 score',r2_score(y_test,y_pred))  
print('MAE',mean_absolute_error(y_test,y_pred))  
R2 score 0.7912364697806901  
MAE 0.2340174549139672
```

```
In [87]: #Gradient Boost  
step1 = ColumnTransformer(transformers=[  
    ('col_tnf',OneHotEncoder(sparse=False,drop='first')), [0,1,7,10,11])  
]),remainder='passthrough')  
step2 = GradientBoostingRegressor(n_estimators=500)  
pipe = Pipeline([  
    ('step1',step1),  
    ('step2',step2)  
])  
pipe.fit(X_train,y_train)  
y_pred = pipe.predict(X_test)  
print('R2 score',r2_score(y_test,y_pred))  
print('MAE',mean_absolute_error(y_test,y_pred))  
R2 score 0.8828824148178134  
MAE 0.15908268928800107
```

```
In [88]: #XGBBoost  
step1 = ColumnTransformer(transformers=[  
    ('col_tnf',OneHotEncoder(sparse=False,drop='first')), [0,1,7,10,11])  
]),remainder='passthrough')  
step2 = XGBRegressor(n_estimators=45,max_depth=5,learning_rate=0.5)  
pipe = Pipeline([  
    ('step1',step1),  
    ('step2',step2)  
])  
pipe.fit(X_train,y_train)  
y_pred = pipe.predict(X_test)  
print('R2 score',r2_score(y_test,y_pred))  
print('MAE',mean_absolute_error(y_test,y_pred))  
R2 score 0.881177343850243  
MAE 0.14496203512600974
```

```
In [89]: #Voting Regressor  
from sklearn.ensemble import VotingRegressor,StackingRegressor  
step1 = ColumnTransformer(transformers=[  
    ('col_tnf',OneHotEncoder(sparse=False,drop='first')), [0,1,7,10,11])  
]),remainder='passthrough')  
rf = RandomForestRegressor(n_estimators=350,random_state=3,max_samples=0.5,max_features=0.75,max_depth=15,bootstrap=True)  
gbrt = GradientBoostingRegressor(n_estimators=100,max_features=0.5)  
xgb = XGBRegressor(n_estimators=25,learning_rate=0.3,max_depth=5)  
et = ExtraTreesRegressor(n_estimators=100,random_state=3,max_samples=0.5,max_features=0.75,max_depth=10,bootstrap=True)  
step2 = VotingRegressor([('rf', rf), ('gbrt', gbrt), ('xgb',xgb), ('et',et)],weights=[5,1,1,1])  
pipe = Pipeline([  
    ('step1',step1),  
    ('step2',step2)  
])  
pipe.fit(X_train,y_train)  
y_pred = pipe.predict(X_test)  
print('R2 score',r2_score(y_test,y_pred))  
print('MAE',mean_absolute_error(y_test,y_pred))  
R2 score 0.8903726162838219  
MAE 0.15900605322976288
```

```
In [90]: #Stacking  
from sklearn.ensemble import VotingRegressor,StackingRegressor  
step1 = ColumnTransformer(transformers=[  
    ('col_tnf',OneHotEncoder(sparse=False,drop='first')), [0,1,7,10,11])  
]),remainder='passthrough')  
estimators = [  
    ('rf', RandomForestRegressor(n_estimators=350,random_state=3,max_samples=0.5,max_features=0.75,max_depth=15,bootstrap=True)),  
    ('gbrt', GradientBoostingRegressor(n_estimators=100,max_features=0.5)),  
    ('xgb', XGBRegressor(n_estimators=25,learning_rate=0.3,max_depth=5))  
]  
step2 = StackingRegressor(estimators=estimators, final_estimator=Ridge(alpha=100))  
pipe = Pipeline([  
    ('step1',step1),  
    ('step2',step2)  
])  
pipe.fit(X_train,y_train)  
y_pred = pipe.predict(X_test)  
print('R2 score',r2_score(y_test,y_pred))  
print('MAE',mean_absolute_error(y_test,y_pred))  
R2 score 0.8795999819487367  
MAE 0.1675181693362485
```

```
In [91]: #Exporting the Model  
import pickle  
pickle.dump(df,open('df.pkl','wb'))  
pickle.dump(pipe,open('pipe.pkl','wb'))
```